High-order cognition is supported by complex but compressible brain activity patterns

Lucy L. W. Owen[1], 2, and Jeremy R. Manning[1,*]
[1]Department of Psychological and Brain Sciences, Dartmouth College, Hanover, NH
[2]Carney Institute for Brain Sciences, Brown University, Providence, RI
[*]Address correspondence to jeremy.r.manning@dartmouth.edu

February 20, 2023

Abstract

We applied dimensionality reduction algorithms to the activity patterns in each experimental condition. Specifically, we sought to understand the "dimensionality" of the neural patterns that were sufficient to decode participants' listening times (or approach was similar to that of Mack et al. 2017). We found that even low-dimensional embeddings of the data were sufficient to accurately decode listening times from the intact story recording, whereas finer temporal scramblings of the story required higher-dimensional embeddings of the data to reach peak decoding accuracy.

Introduction

Large-scale networks, including the human brain, may be conceptualized as occupying one or more positions along on a continuum. At one extreme, every node is fully independent of every other node. At the other extreme, all nodes are fully coupled and behave identically. Each extreme optimizes key properties of how the network functions. When every node is independent, the network is maximally *expressive*: if we define the network's "state" as the total set of activity patterns across nodes, then every state is equally reachable by a network with fully independent nodes. On the other hand, a fully coupled network optimizes *robustness*: any subset of nodes, other than the entire network, may be removed from the network without any loss of function or expressive power. Note that a given set of nodes might reconfigure its connections or behaviors under different circumstances to change its position along this continuum according to the needs at hand. Presumably, most systems tend to occupy positions between the above extremes. We wondered: might the human brain reconfigure itself to be more flexible or more robust according to ongoing demands?

We're interested in the complexity of brain patterns that underly different types of thoughts. To explore this question space, we will take brain patterns recorded under different experimental conditions used in Aim 2, and project them into lower dimensional spaces using principle components analysis. We can then ask how well those low-dimensional embeddings of the data retain cognitively relevant information like when in a story someone is listening to.

This work has been inspired, in part, by Mack et al. (2020). In this paper, they investigated the role of the prefrontal cortex in filtering out irrelevant content. Specifically, they looked at if the vmPFC performs data reduction on incoming information through compression. This was motivated, in part, by orbital frontal cortex (OFC) compression in rats (Zhou et al., 2019). They studied this using a learning paradigm in which participants had to classify insects based on different numbers of feature dimensions. The idea was that participants in some learning blocks, participants could identify the insects based on one feature (low complexity) or several features (high complexity), but importantly the stimuli remained the same across all learning problems. They found that complexity and compression had an inverse relationship; the lower complexity of a conceptual space, the higher the degree of compression. Building on this idea, we wonder if varying degrees of compression is performed throughout the brain. We also want to test this idea, but using varying levels of engagement listening to a naturalistic stimuli.

To understand the degree of compression throughout the brain during cognition, we will use the same fMRI data from Aim 2, collected while participants listened to a story in different scrambling conditions. We will measure the degree that multivoxel activation patterns are compressed during story listening using principle components analysis (PCA) a method for low-rank approximation of multidimensional data (Eckart & Young, 1936). We will explore this using decoding accuracy as a function of the number of components, or dimensions, in the low-dimensional space under different cognitive conditions.

You can imagine two reasonable predictions of how cognition is reflected in brain patterns. The first is as our thoughts become more complex, they are supported by more complex brain patterns, and require more components to decode. The second is that when thoughts are deeper and more complicated, the units of neural activity would carry more information, and would require therefore fewer components to decode.

This idea can be explored in this visual analogy (Fig. **??**) for neural compression. Here there are two images of pies, the top pie is more complex than the bottom. On the left we're illustrating that it takes fewer components to reach the same 95 percent variance explained in the less complex pie, which corresponds to higher compression. However, on the right with very few components similar variance is explaining both pies.

We investigated the dimensionality of neural patterns by training classifiers using more and more principle components. Or, in other words, we used less and less compression to decode. We applied the approach to a neuroimaging dataset comprising data collected as participants listened to a story varying in cognitive richness (Simony et al., 2016).

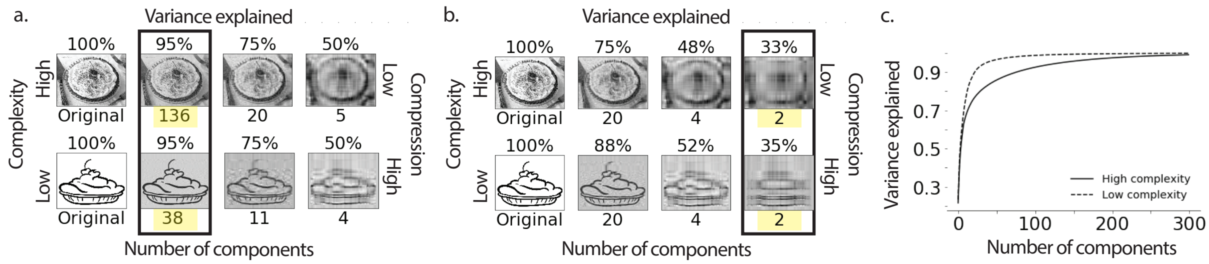

Figure 1: **Illustration of compression.** Visual analogy for neural compression. Here are 2 images of pies, one more complex than the other. **a.** It takes fewer components to reach the same percent variance explained in the less complex pie, which corresponds to higher compression. **b.** However, with very few components, similar variance is explained in both pies. **c.** Plots the cumulative explained variance for more and more components.

Evaluation metrics

We will evaluate the degree of compression of held-out neuroimaging data by assessing the time at which it was collected. We will use this evaluation (timepoint decoding) as a proxy for gauging how much explanatory power the compressed data held with respect to the observed data.

Timepoint decoding

To explore how compression varies with complexity, we will use a previous neuroimaging dataset Simony et al. (2016) in which participants listened to an audio recording of a story; 36 particpants listen to an intact version of the story, 17 participants listen to time-scrambled recordings of the same story where paragraphs were scrambled, 36 particpants listen to word-scrambled version and 36 participants lay in rest condition.

Following the analyses conducted by (HTFA) Manning et al. (2018), we first apply *hierarchical topographic factor analysis* (HTFA) to the fMRI datasets to obtain a time series of 700 node activities for every participant. We then apply dimensionality reduction (Incremental PCA) for each group.

We then compare the groups' activity patterns (using Pearson correlations) to estimate the story times each corresponding pattern using more and more principle components.

To assess decoding accuracy, we randomly divide participants for each stimulus into training and testing groups. We then compare the groups' activity patterns (using Pearson correlations) to estimate the story times each corresponding pattern using more and more principle components (as the data became less compressed). Specifically, we ask, for each timepoint: what are the correlations between the first group's and second group's activity patterns at each order. We note that the decoding test we used is a conservative in which we count a timepoint label as incorrect if it is not an exact match.

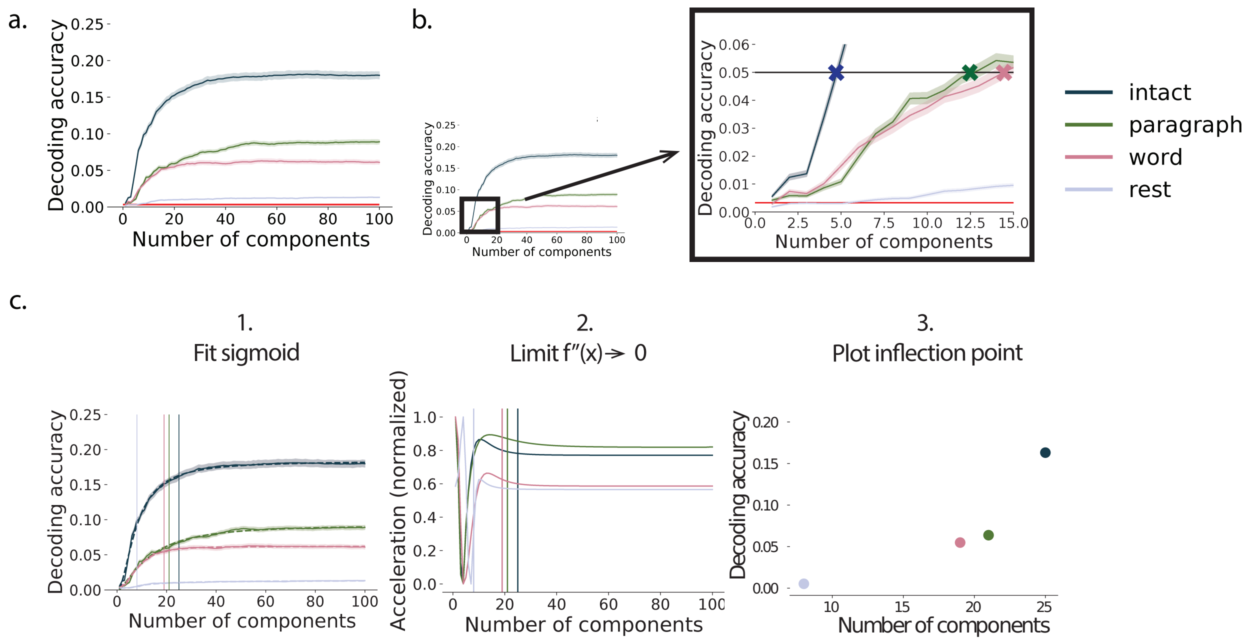

Figure 2: **Decoding accuracy. a. Decoding accuracy by number of components.** Ribbons of each color display cross-validated decoding performance for each condition (intact, paragraph, word, and rest). Decoders were trained using increasingly more principle components and displayed relative to chance (red line). **b. Fixed decoding accuracy by number of components.** We zoom in on the plot shown in **a.** and add a line denoting fixed decoding accuracy (.05). We plot where the intact, paragraph, and word conditions intersect. **c. Explanation of inflection metric.** First the we fit a sigmoid function to the decoding accuracy by number of components. Second, we found where the second derivative is both positive and less than .0001. Last, we then plot that inflection point as a single metric to capture the slope and asymptote of the curve.

Results

By training classifiers using more and more principle components to decode, and comparing across conditions with varying degrees of cognitive richness, we can assess the explanatory power of the compressed data held with respect to the observed data (see *Methods*). We note that our primary goal was not to achieve perfect decoding accuracy, but rather to use decoding accuracy as a benchmark for assessing whether different neural features specifically capture cognitively relevant brain patterns.

Prior work has shown participants share similar neural responses to richly structured stimuli when compared to stimuli with less structure Simony et al. (2016). We replicate this finding, showing as complexity of the stimulus increases, decoding accuracy increases (Fig. 2, a.). Additionally, we found that as complexity of the stimuli increases, we need fewer components to decode the same amount (Fig. 2, b.). However, we also found that as complexity of the stimuli increases, more components are required to reach peak decoding accuracy (Fig. 2, c.). We posit that as the complexity of our thoughts increases, neural compression decreases. However, as our thoughts become deeper and richer, more reliable information is available at higher neural compression.

We also wondered how this compression would change across brain regions. We repeated the analysis but limited the brain hubs to 7 networks using the Yeo et al. (2011) network parcellation shown here in the inflated brain (Fig. 3, d.). We found that as complexity of the stimuli increases, decoding accuracy increases with higher cognitive areas. (Fig. 3).

We were also curious how compression would change across time. If, there is some understanding of the narrative that accumulates over time, we should be able to see that difference. We found increases in decoding accuracy with the same number or fewer components for more complex, cognitively rich, conditions. We also found decreases in decoding accuracy for the word-scrambled and rest condition.

Overall, we found that as story listening conditions become more complex, more components are required to decode. We also found we could decode better with more impoverished data when there is the underlying structure of the narrative providing more cognitive richness. We posit that as the complexity of our thoughts increases, neural compression decreases. However, as our thoughts become deeper and richer, more reliable information is available at higher neural compression.

Discussion

- We trained classifiers using more and more principle components to decode, and compared across conditions with varying degrees of cognitive richness. -We found that as listening conditions become more cognitively rich, decoding accuracy increased. -Also, decoding accuracy increased as understanding of the narrative accumulated over time, in more complex listening conditions. - Decoding accuracy also increased in higher cognitive areas, in more complex listening conditions. -We found that as story listening conditions become more complex, more components are required to decode. -We also found we could decode better with more impoverished data when there is the underlying structure of the narrative providing more cognitive richness. -We posit that as the complexity of our thoughts increases, neural compression decreases. However, as our thoughts become deeper and richer, more reliable information is available at higher neural compression.

Based on prior work (**?**) and following the direction of the field (Turk-Browne, 2013) we think our thoughts might be encoded in dynamic network patterns, and possibly higher order network patterns (Fig. **??**). We sought to test this hypothesis by developing an approach to inferring high-order network dynamics from timeseries data.

One challenge in studying dynamic interactions is the computational resources required to calculate higher-order correlations. We developed a computationally tractable model of network dynamics (Fig. **??**)

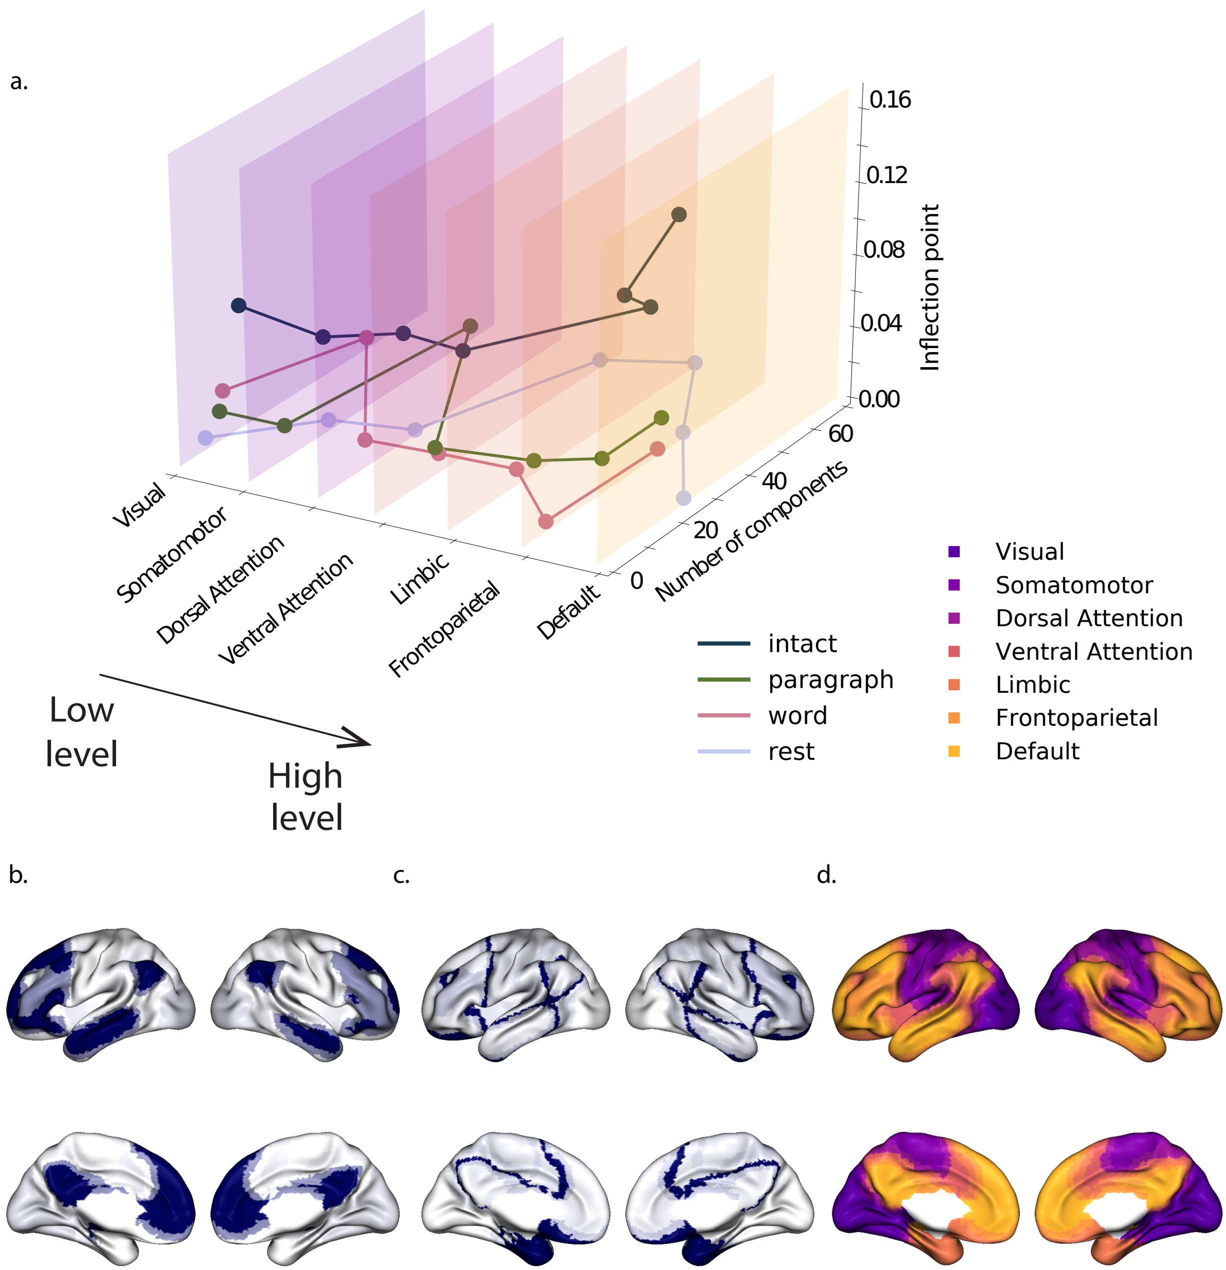

Figure 3: **Inflection points by network.** **a.** Inflection point was calculated as explained in Fig. 2, b. Analyses were limited by the brain networks (using the Yeo et al. (2011) network parcellation) and arranged in increasing order relative to the intact condition. **b. and c.** For the total time in the intact condition, we are plotting the relative inflection points (**b.**) and corresponding number of components (**c.**) by network. **d.** The network parcellation defined by Yeo et al. (2011) is displayed on the inflated brain maps. The colors and network labels serve as a legend for **a.** and **d.**

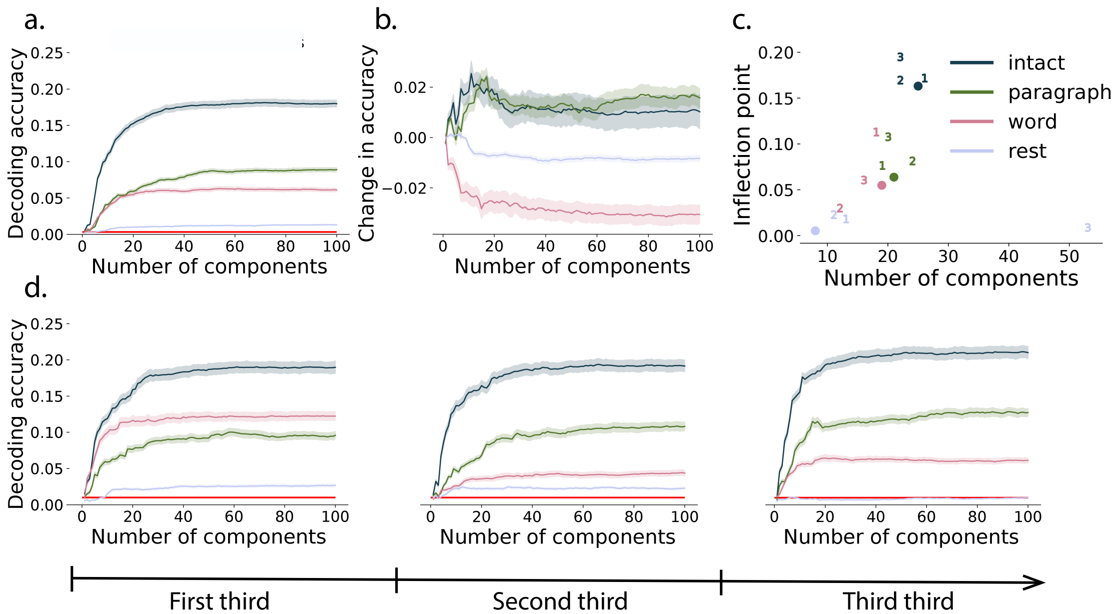

Figure 4: **Inflection points by thirds.** **a.** Decoding accuracy by number of components not broken into thirds (Fig. 2 a.). **b. and c.** Quantifying changes in decoding accuracy across time. **b.** Slope of decoding accuracy was calculated by fitting a regression line for each component/condition for each third. **c.** We also repeated the analysis (Fig. 2, b.) to obtain the inflection point for each condition and for each third. **d.** Decoding accuracy by number of components for each third of the scan time. We repeated the same analysis in Fig. 2 a. but breaking the scan time for each condition into 3 intervals.

that takes in a feature timeseries and outputs approximated first-order dynamics (i.e., dynamic functional correlations), second-order dynamics (reflecting homologous networks that dynamically form and disperse), and higher-order network dynamics (up to tenth-order dynamic correlations).

We first validated our model using synthetic data, and explored how recovery varied with different underlying data structures and kernels. We then applied the approach to an fMRI dataset (Simony et al., 2016) in which participants listened to an audio recording of a story, as well as scrambled versions of the same story (where the scrambling was applied at different temporal scales). We trained classifiers to take the output of the model and decode the timepoint in the story (or scrambled story) that the participants were listening to. We found that, during the intact listening condition in the experiment, classifiers that incorporated higher-order correlations yielded consistently higher accuracy than classifiers trained only on lower-order patterns (Fig. **??**, a.&d.). By contrast, these higher-order correlations were not necessary to support decoding the other listening conditions and (minimally above chance) during a control rest condition. This suggests that the cognitive processing that supported the most cogntively rich listening conditions involved second-order (or higher) network dynamics.

Although we found decoding accuracy was best when incorporating higher-order network dynamics for all but rest condition, it is unclear if this is a product of the brain or the data collection technique. It could be that the brain is second-order or it could be that fMRI can only reliably give second-order interactions. Exploring this method with other data collection technique will be important to disentangle this question.

Concluding remarks

How can we better understand how brain patterns change over time? How can we quantify the potential network dynamics that might be driving these changes? One way to judge the techniques of the future is to look at the trajectory of the fMRI field so far has taken so far (Fig. **??**). The field started with univariate activation, measuring the average activity for each voxel. Analyses of multivariate activation followed, looking at spatial patterns of activity over voxels. Next, correlations of activity were explored, first with measures like resting connectivity that take temporal correlation between a seed voxel and all other voxels then with full connectivty that measure all pairwise correlations. Additionally, this path of increasing complexity also moved from static to dynamic measurements. One logical next step in this trajectory would be dynamic higher-order correlations. We have created a method to support these calculations by scalably approximating dynamic higher-order correlations.

Acknowledgements

We acknowledge discussions with Luke Chang, Hany Farid, Paxton Fitzpatrick, Andrew Heusser, Eshin Jolly, Qiang Liu, Matthijs van der Meer, Judith Mildner, Gina Notaro, Stephen Satterthwaite, Emily Whitaker, Weizhen Xie, and Kirsten Ziman. Our work was supported in part by NSF EPSCoR Award Number 1632738 to J.R.M. and by a sub-award of DARPA RAM Cooperative Agreement N66001-14-2-4-032 to J.R.M. The content is solely the responsibility of the authors and does not necessarily represent the official views of our supporting organizations.

Author contributions

Concept: J.R.M. and L.L.W.O. Implementation: L.L.W.O., and J.R.M. Analyses: L.L.W.O and J.R.M.

References

Eckart, C., & Young, G. (1936). The approximation of one matrix by another of lower rank. *Psychometrika*, *1*, 211–218.

Mack, M. L., Preston, A. R., & Love, B. C. (2020). Ventromedial prefrontal cortex compression during concept learning. *Nature Communications*, *11*(1), 46. Retrieved from `https://doi.org/10.1038/s41467-019-13930-8` doi: 10.1038/s41467-019-13930-8

Manning, J. R., Zhu, X., Willke, T. L., Ranganath, R., Stachenfeld, K., Hasson, U., . . . Norman, K. A. (2018). A probabilistic approach to discovering dynamic full-brain functional connectivity patterns. *NeuroImage*, *180*, 243–252.

Simony, E., Honey, C. J., Chen, J., & Hasson, U. (2016). Dynamic reconfiguration of the default mode network during narrative comprehension. *Nature Communications*, *7*(12141), 1–13.

Turk-Browne, N. B. (2013). Functional interactions as big data in the human brain. *Science*, *342*, 580–584.

Yeo, B. T. T., Krienen, F. M., Sepulcre, J., Sabuncu, M. R., Lashkari, D., Hollinshead, M., . . . Buckner, R. L. (2011). The organization of the human cerebral cortex estimated by intrinsic functional connectivity. *Journal of Neurophysiology*, *106*(3), 1125–1165.

Zhou, J., Gardner, M. P. H., Stalnaker, T. A., Ramus, S. J., Wikenheiser, A. M., Niv, Y., & Schoenbaum, G. (2019). Rat orbitofrontal ensemble activity contains multiplexed but dissociable representations of value and task structure in an odor sequence task. *Current Biology*, *29*(6), 897–907.e3. Retrieved from `https://www.sciencedirect.com/science/article/pii/S0960982219300776` doi: https://doi.org/10.1016/j.cub.2019.01.048